

Scalable Cloud Architectures

Dr Dan Schien – COMSM0010
Lecture 11

bristol.ac.uk



Review – COMSM0010 so far...

Security

Hadoop

Spark

DevOps

Scalable Software Architecture

Distributed
File Systems

Databases
(*SQL)

Databases
(Graph)

Stream
Processing

Legal

*aaS

VM

Containers

Orchestration

Serverless

Economics

Introduction

Coursework

What is Cloud Native?

- Container based
 - Capabilities
- Elastic
- Microservices
 - Polyglot
- DevOps





DATA CENTRE

SOFTWARE

SECURITY

DEVOPS

BUSINESS

PERSONAL TECH

SCIENCE

EN



everywhere
you want to be

A secure way
to pay abroad.

Data Centre ▶ Cloud

Yale Security Fail: 'Unexpected load' caused systems to crash, whacked our Smart Living Home app

All working now says biz. No, no, no, no, say
customers, it is NOT!

By [Paul Kunert](#) 19 Oct 2018 at 16:30

59 SHARE ▼



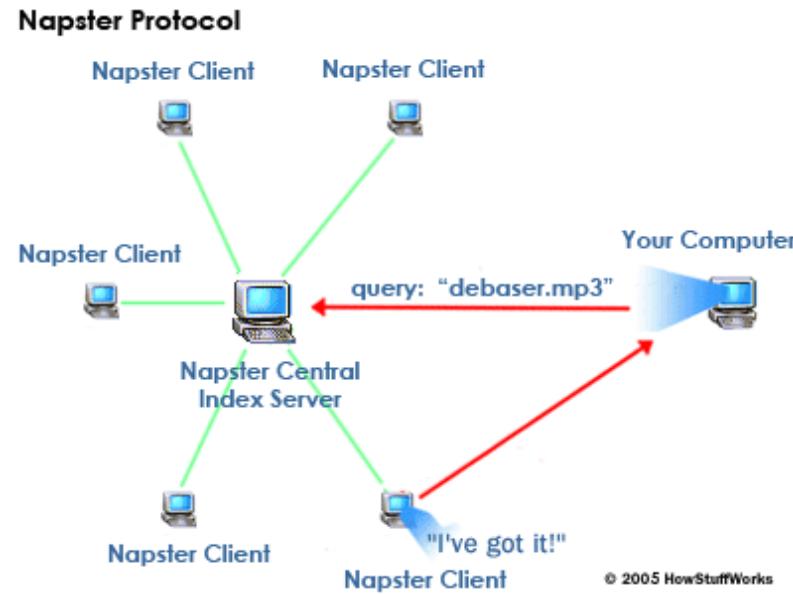
Software Architectures

Definition Software Architecture

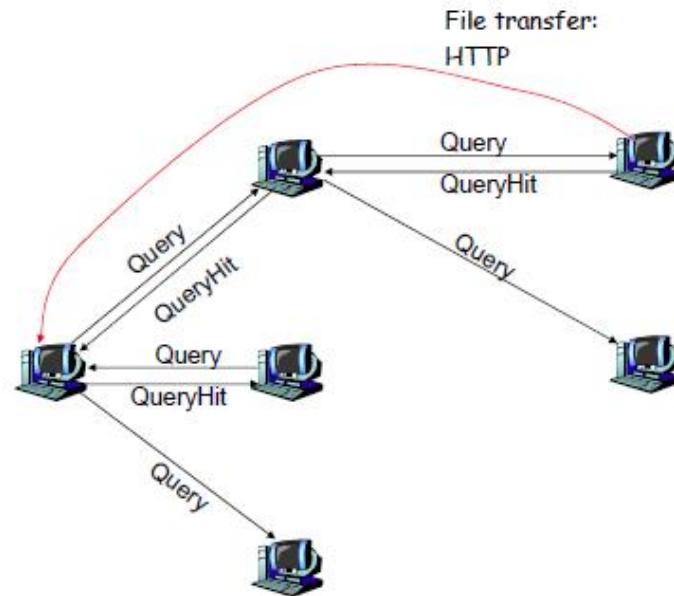
- Software architecture is a set of **structures** needed to **reason** about a system. It includes software components, the **relations** between them and properties of both.
- Architecture might be implicit
 - If you're building any significant system. You cannot not have an architecture. You might just not been have documented or communicated explicitly.

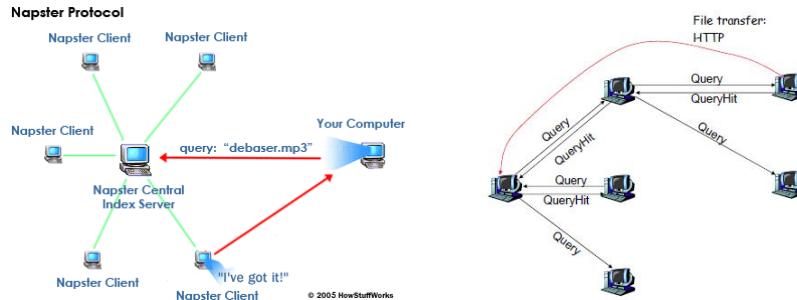
**What is the different between
Napster and Gnutella?**

Napster



Gnutella



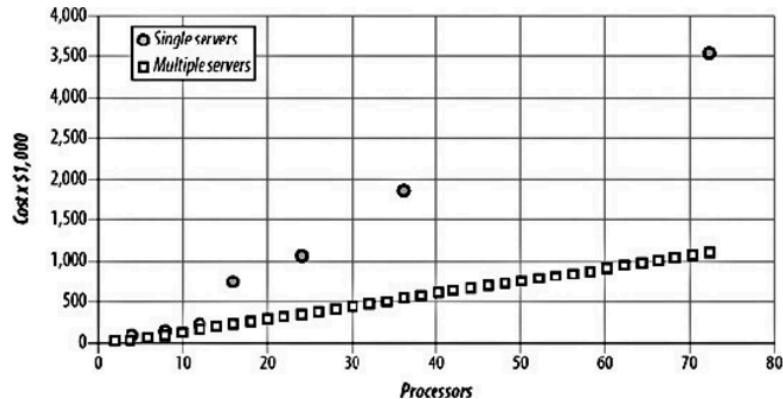


	Napster	Gnutella
Architecture	Centralised Index	Unstructured Using simple flooding
Advantages	+ Highly efficient data lookup + Rapidly adapts to changes in network	+ Entirely decentralized, pure P2P network + Highly resistant to failure
Disadvantages	- Questionable scalability - Vulnerable to censorship, failure, attack	- Search is time-consuming - Network typically scales poorly

Principles of Scaling Systems

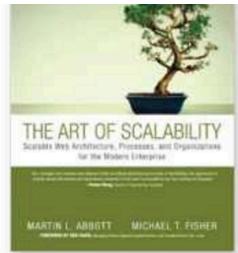
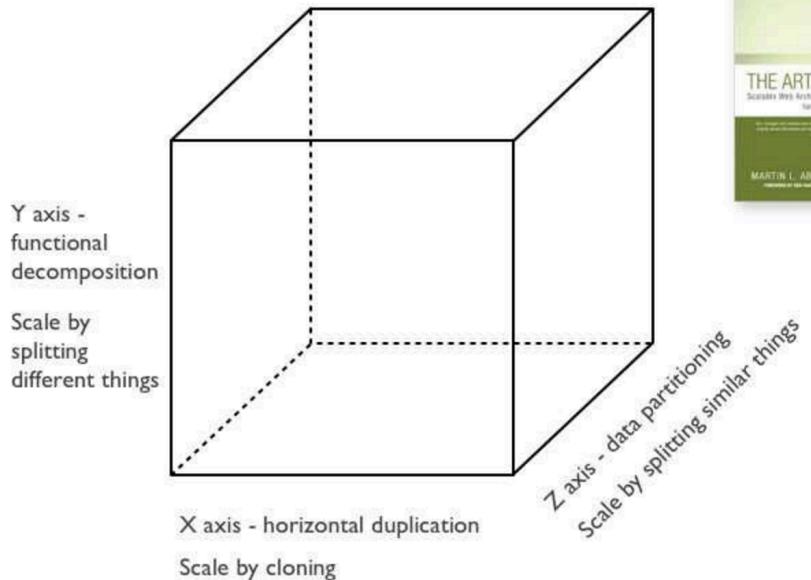
Horizontal vs Vertical Scaling

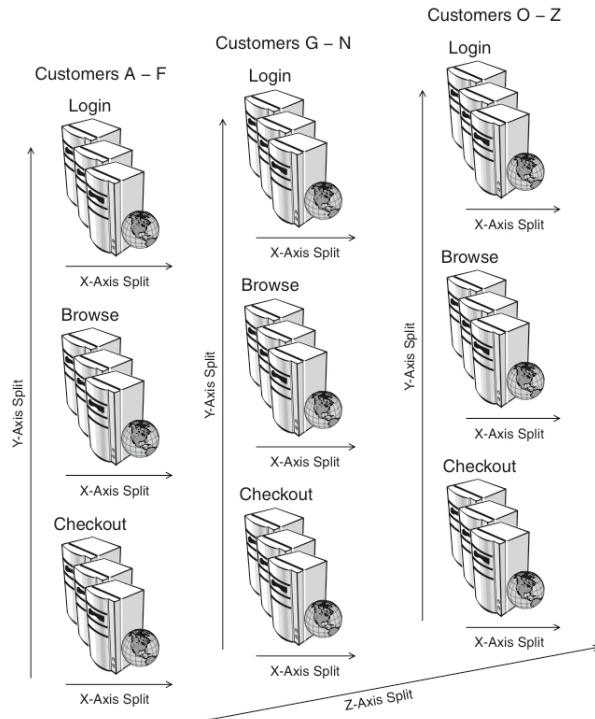
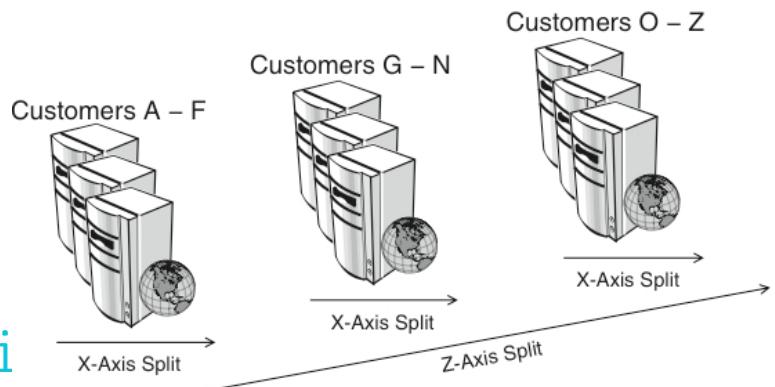
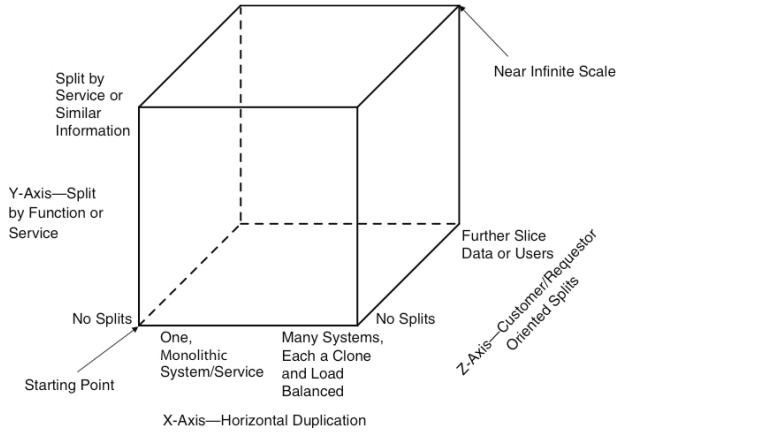
- Vertical (Scale Up) bigger machines
 - no need to change your architecture
- Horizontal (Scale Out) more machines
- Trade-off: Cost of metal vs cost of developers



<http://www.amazon.com/dp/0596102356>

The Scale Cube (of Scaling Out – horizontally only)

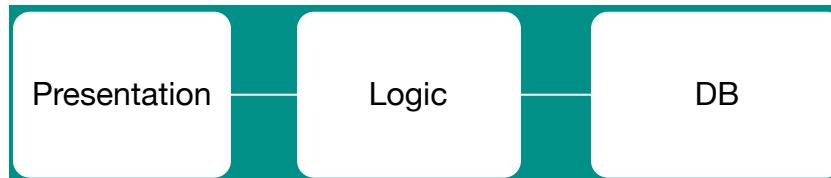




- X axis (horizontal duplication) unbiased cloning of services and data
 - each clone can do the work of all other clones. work is distributed among clones without bias
 - inefficient compared to alternatives
 - easy to do
- Y axis (split by function or service) refers to isolating and making scalable individual responsibility of components
 - needs to be split in the code base
 - more costly than x-axis
- Z axis - partitioning the domain of incoming requests
 - data partitioning, split relative to client
 - improves fault tolerance, cache performance
 - more costly than other two

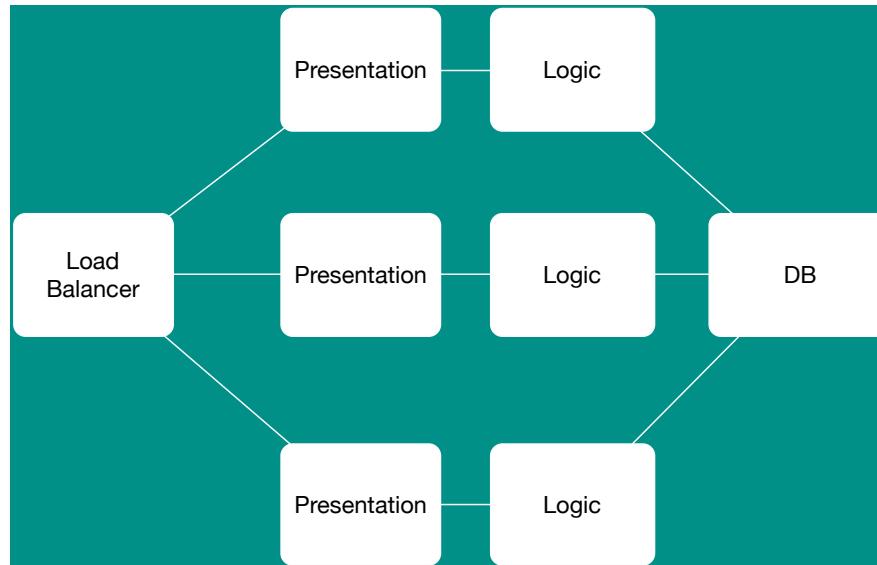
How to scale out? Layered Architectures

- Consider this simple application, might be a teaching and learning system – such as Blackboard
- Presentation layer provides a mobile API and web UI
- Logic layer has the business logic, book hotels and flights, etc

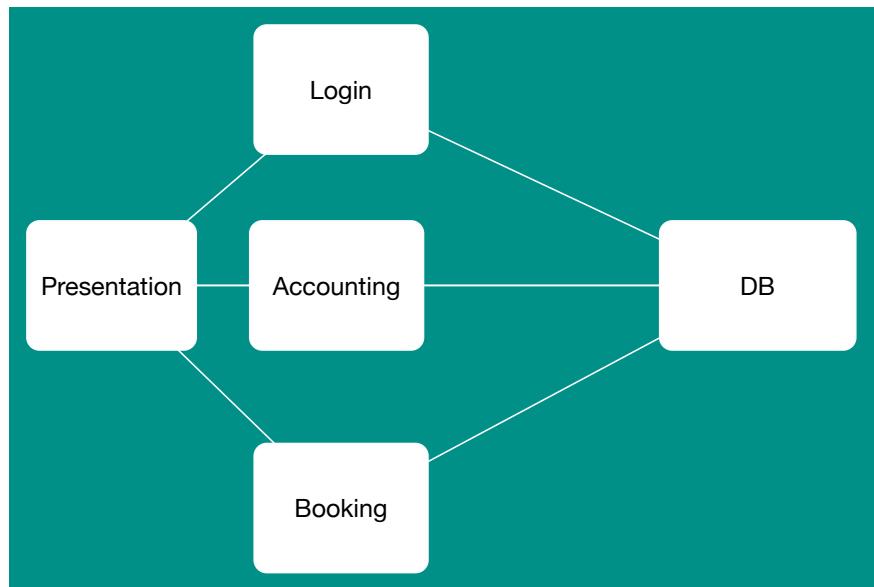


X-Axis scaling

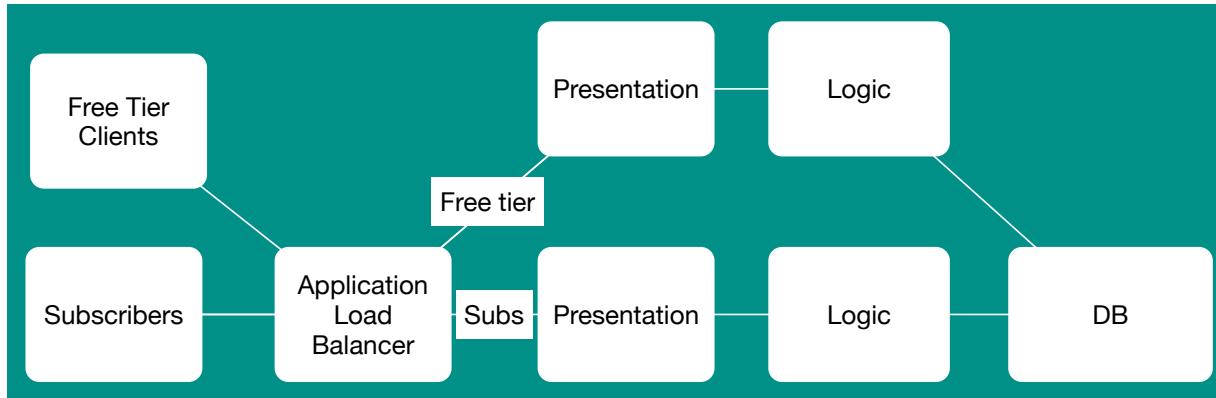
Scale application logic or presentation layers



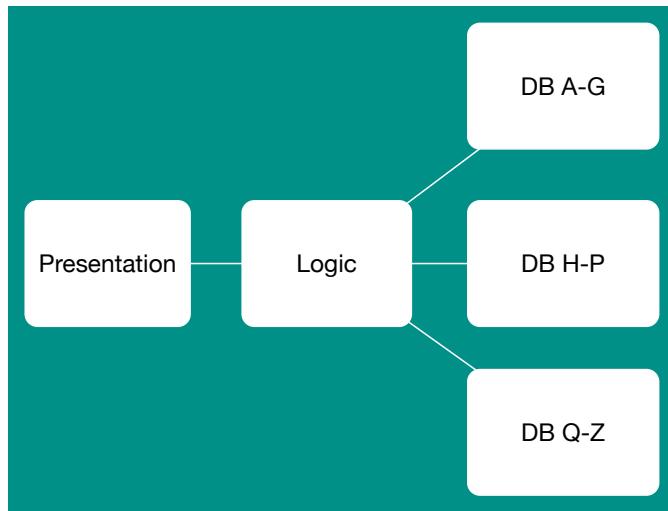
Y-Axis: decompose by function



Z-Axis scaling 1: partition per request type



Z-Axis 2: Scaling by partitioning the data



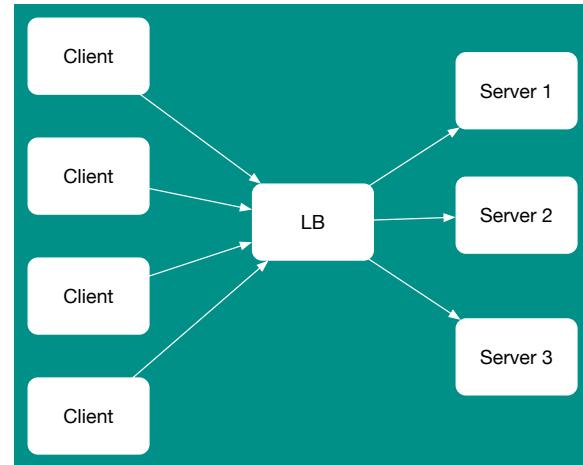
Architectural components and patterns for scalable systems

Decoupled Components

- Independent scalability of components
- Preliminary for Fault Tolerance and High Availability
- Essential requirement for the use of these principles: APIs defined for all components (e.g. HTTP REST, RPC or typed messages)
- Four important mechanisms to **decouple**
 - Load Balancers
 - Message Queues
 - Message Topics
 - Service Registries

Load Balancers

- distributing requests
- manage availability -> HA
- Perform health check
- Session affinity / Sticky Sessions
- list of backend servers
- load balancing policy
- SSL
- Alternative: DNS or Router based



Load Balancers

- Hides the server from client access

I

[Airbnb](#)^[11]
[Amazon.com](#)^[8]
[Ancestry.com](#)^{[12][13]}
[The A.V. Club](#)^[14]
[BBC](#)^[13]
[The Boston Globe](#)^[11]
[Box](#)^[15]
[Business Insider](#)^[13]
[CNN](#)^[13]
[Comcast](#)^[16]
[CrunchBase](#)^[13]
[DirecTV](#)^[13]
[The Elder Scrolls Online](#)^{[13][17]}
[Electronic Arts](#)^[16]
[Etsy](#)^{[11][18]}
[FiveThirtyEight](#)^[13]
[Fox News](#)^[19]
[The Guardian](#)^[19]
[GitHub](#)^{[11][16]}
[Grubhub](#)^[20]

[HBO](#)^[13]
[Heroku](#)^[21]
[HostGator](#)^[13]
[iHeartRadio](#)^{[12][22]}
[Imgur](#)^[23]
[Indiegogo](#)^[12]
[Mashable](#)^[24]
[National Hockey League](#)^[13]
[Netflix](#)^{[13][19]}
[The New York Times](#)^{[11][16]}
[Overstock.com](#)^[13]
[PayPal](#)^[18]
[Pinterest](#)^{[16][18]}
[Pixlr](#)^[13]
[PlayStation Network](#)^[16]
[Qualtrics](#)^[12]
[Quora](#)^[13]
[Reddit](#)^{[12][16][18]}
[Roblox](#)^[25]
[Ruby Lane](#)^[13]
[RuneScape](#)^[12]

[SaneBox](#)^[21]
[Seamless](#)^[23]
[Second Life](#)^[26]
[Shopify](#)^[11]
[Slack](#)^[23]
[SoundCloud](#)^{[11][18]}
[Squarespace](#)^[13]
[Spotify](#)^{[12][16][18]}
[Starbucks](#)^{[12][22]}
[Storify](#)^[15]
[Swedish Civil Contingencies Agency](#)^[27]
[Swedish Government](#)^[27]
[Tumblr](#)^{[12][16]}
[Twilio](#)^{[12][13]}
[Twitter](#)^{[11][12][16][18]}
[Verizon Communications](#)^[16]
[Visa](#)^[28]
[Vox Media](#)^[29]
[Walgreens](#)^[13]
[The Wall Street Journal](#)^[19]

1	// root	xc3511	// root	vizxv	// root	admin
2	// admin	admin	// root	888888	// root	xmhdpic
3	// root	default	// root	juantech	// root	123456
4	// root	54321	// support	support	// root	(none)
5	// admin	password	// root	root	// root	12345
6	// user	user	// admin	(none)	// root	pass
7	// admin	admin1234	// root	1111	// admin	smcadmin
8	// admin	1111	// root	666666	// root	password
9	// root	1234	// root	klv123	// Administrator	admin
10	// service	service	// supervisor	supervisor	// guest	guest
11	// guest	12345	// guest	12345	// admin1	password
12	// administrator	1234	// 666666	666666	// 888888	888888
13	// ubnt	ubnt	// root	klv1234	// root	Zte521
14	// root	hi3518	// root	jvbzd	// root	anko
15	// root	zlxx.	// root	7ujMko0vizxv	// root	7ujMko0admin
16	// root	system	// root	ikwb	// root	dreambox
17	// root	user	// root	realtek	// root	00000000
18	// admin	1111111	// admin	1234	// admin	12345
19	// admin	54321	// admin	123456	// admin	7ujMko0admin
20	// admin	1234	// admin	pass	// admin	meinsm
21	// tech	tech				

Dyn DoS Attack

- Oct 2016, Much of Eastern US Internet
- Mirai botnet
- IoT
 - 600,000 devices
- Two previous attacks
 - OVH (1.1 Tb/s)
 - Krebs on Security (623 Gb/s)
 - Akamai dropped site

https://en.wikipedia.org/wiki/2016_Dyn_DoS_attack

<https://minecraft.net/en-us/store/>

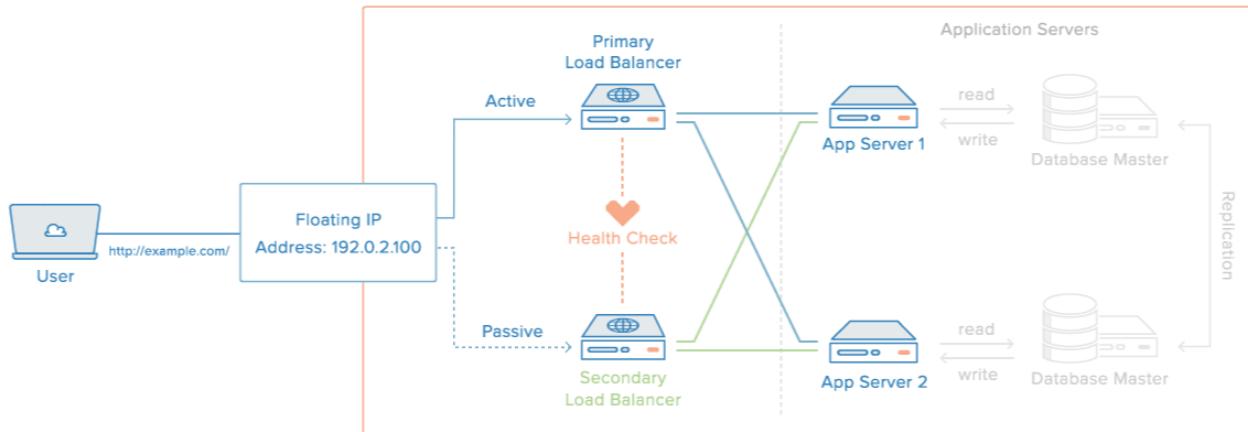
Session affinity / Sticky Sessions

- Sticky LBs
 - always routing a client's requests to the same server instance
 - IP Hash
- Cookies managed by load balancer (duration based)
- Cookies managed by application cookie
- Session Replication - Session-aware clusters
 - Shared session DB or cache
 - Eg. Redis shared
- Session stored with client
 - Exposes application internals

LB Algorithms

- Even Task Distribution / Round Robin
 - Sequential distribution of work among server
 - ignores the difference in the work required
- Weighted Round Robin
- Least Connections

HA LBs



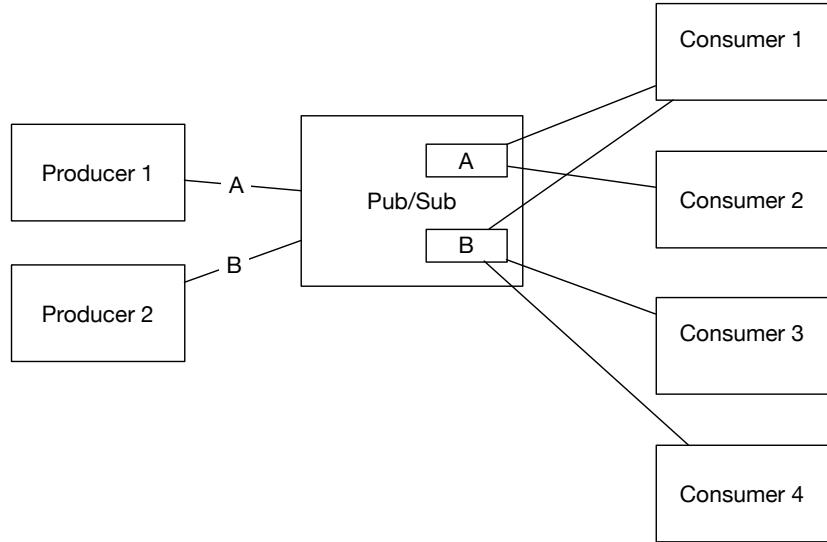
- 1 Active/Passive Cluster is healthy
- 2 Primary node fails
- 3 Floating IP is assigned to Secondary node

🇺🇸 NYC3 DATACENTER

<https://www.digitalocean.com/community/tutorials/an-introduction-to-haproxy-and-load-balancing-concepts>

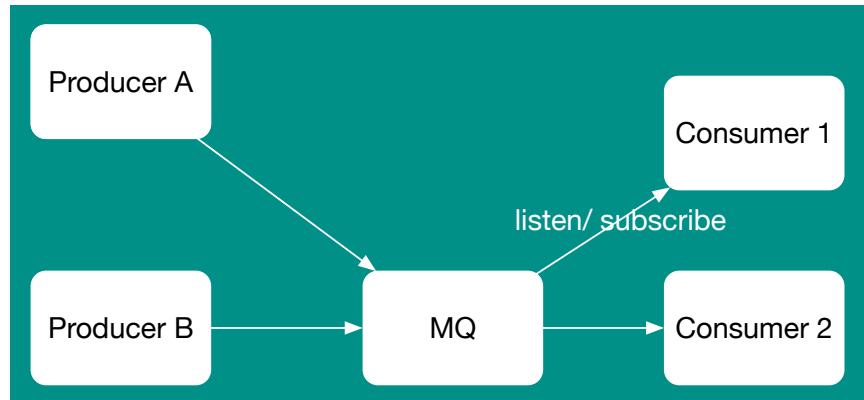
Message Topics

- Messages are immediately pushed to subscribers
- Decouple producers from subscribers
- Concurrent processing
- Independent scalability of components
- Push, One-to-Many distribution

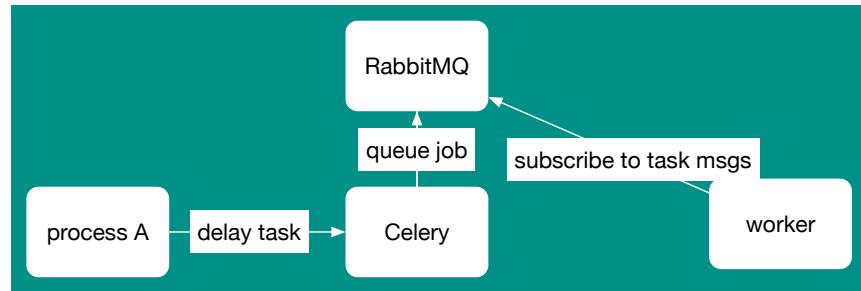


Message Queues

- Asynchronous: Queue it now, run it later. Until a consumer is ready to process
- Decoupling: Separates application logic.
- Scalable: Many workers can process individual jobs in a queue.
- Introduce Latency



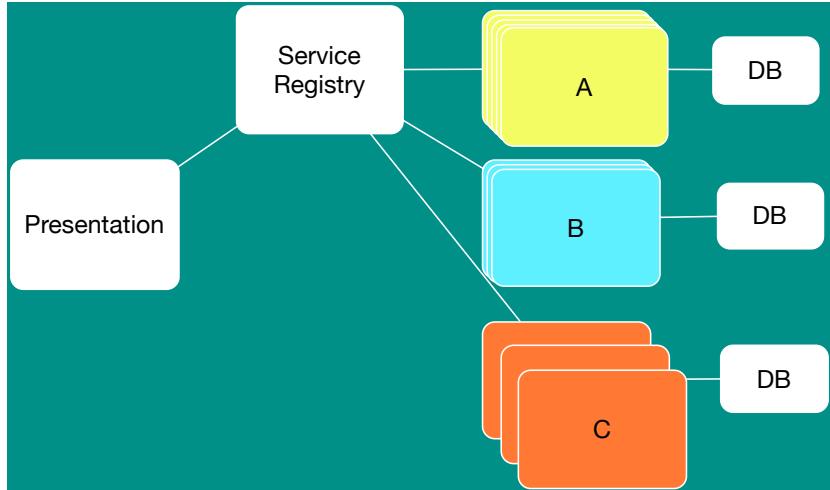
Example: Celery



```
{"id": "4cc7438e-afd4-4f8f-a2f3-f46567e7ca77",  
 "task": "celery.task.PingTask",  
 "args": [],  
 "kwargs": {}, "retries": 0,  
 "eta": "2009-11-17T12:30:56.527191"}
```

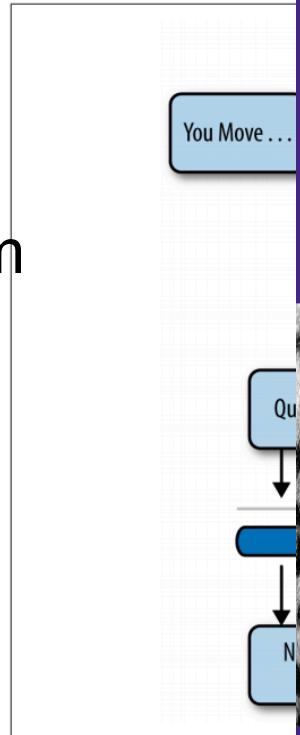
Service Registries

- Resolve Addresses for Names
- Based on HA, transactional data stores
- Examples: Apache Zookeeper, Eureka



Example:

- Customer Processes writes to Pub/Sub system
- Quote Process and Claim Process act concurrently



Software Architecture Patterns



Mark Richards

Automation

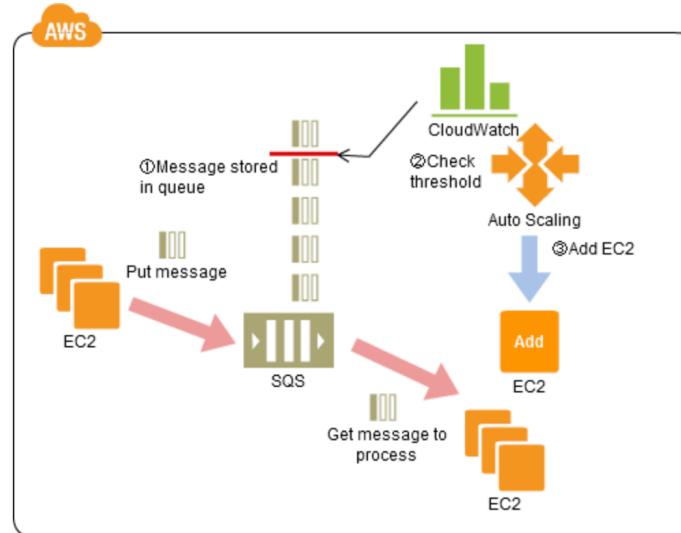
Scaling instances in your Architecture cannot be manual

- Elasticity controller
- Metrics- CPU, Mem, Disk
- When exceeded -> start new nodes
- When underused -> stop nodes

Example: Job Observer Pattern

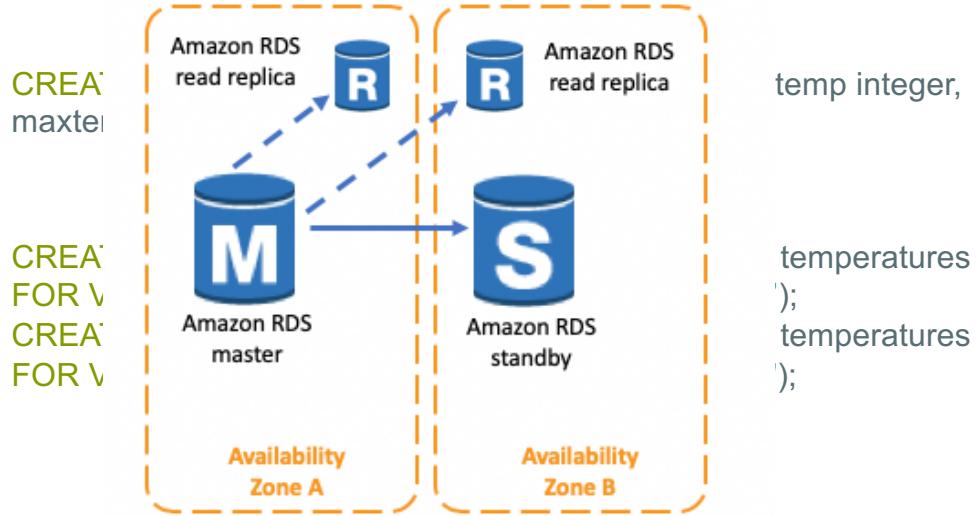
- Queue Scaling
- Watch a job queue and scale depending on that

[http://en.clouddesignpattern.org/index.php/
CDP:Job_Observer_Pattern](http://en.clouddesignpattern.org/index.php/CDP:Job_Observer_Pattern)



Database Scaling

- Read-replicas
- Database sharding
 - PostgreSQL 11 with FDW

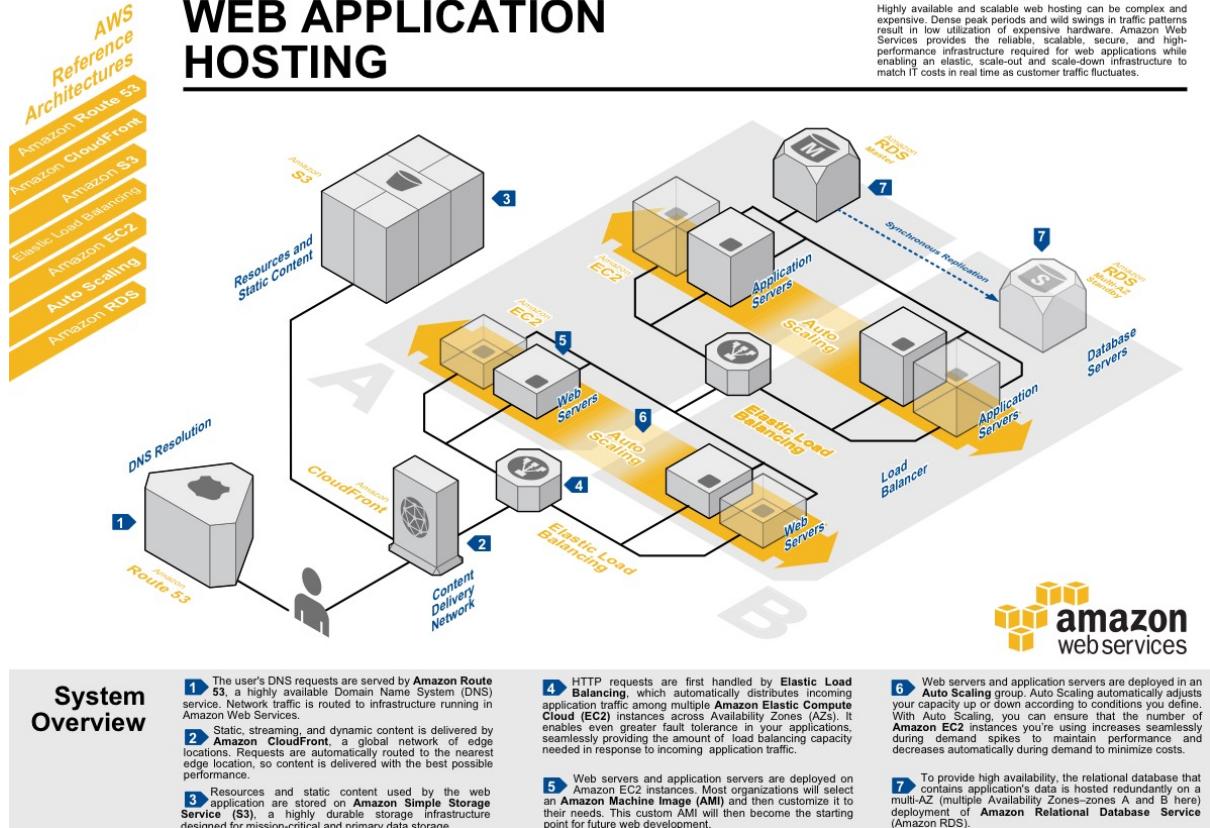


```
CREATE SERVER box2 FOREIGN DATA WRAPPER p  
dbname 'box2db');
```

```
CREATE FOREIGN TABLE temperatures_2016 PARTITION OF temperatures FOR VALUES FROM ('2016-01-  
01') TO ('2017-01-01') SERVER box2;
```

Architectural Patterns

WEB APPLICATION HOSTING



Thursday

- More on Microservices
- Monitoring and Tracing