

Microservices and the CNCF

Dr Dan Schien – COMSM0010
Lecture 12

bristol.ac.uk



Review – COMSM0010 so far...

Review

Security

Hadoop

Spark

DevOps

Scalable Software Architecture

Distributed
File Systems

Databases
(*SQL)

Databases
(Graph)

Stream
Processing

Legal

*aaS

VM

Containers

Orchestration

Serverless

Economics

Introduction

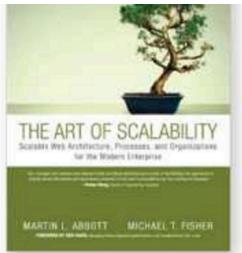
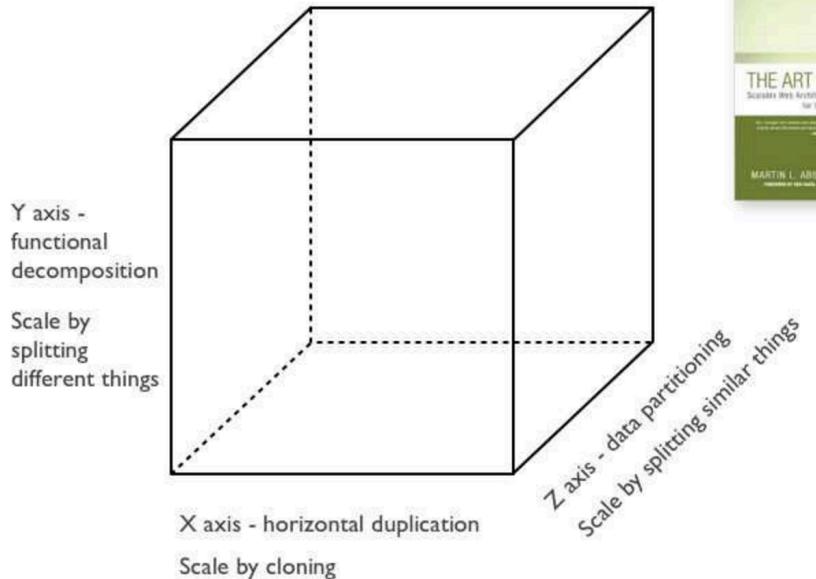
Coursework

Learning outcomes

- Explain the architectural needs that Microservices address
- Describe the organizational goals that Microservices meet
- Understand some of the challenges that Microservices bring
- Recognize the role of the tools from the CNCF landscape
- Describe how observability of Microservices can be addressed

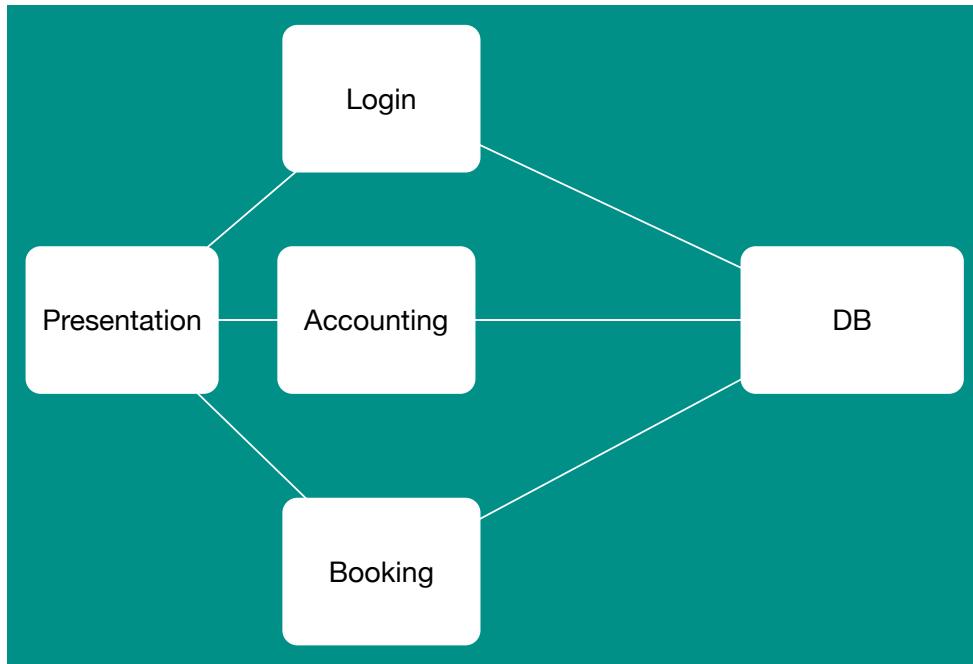
The Scale Cube (of Scaling Out – horizontally only)

Review



Review

Y-Axis: decompose by function



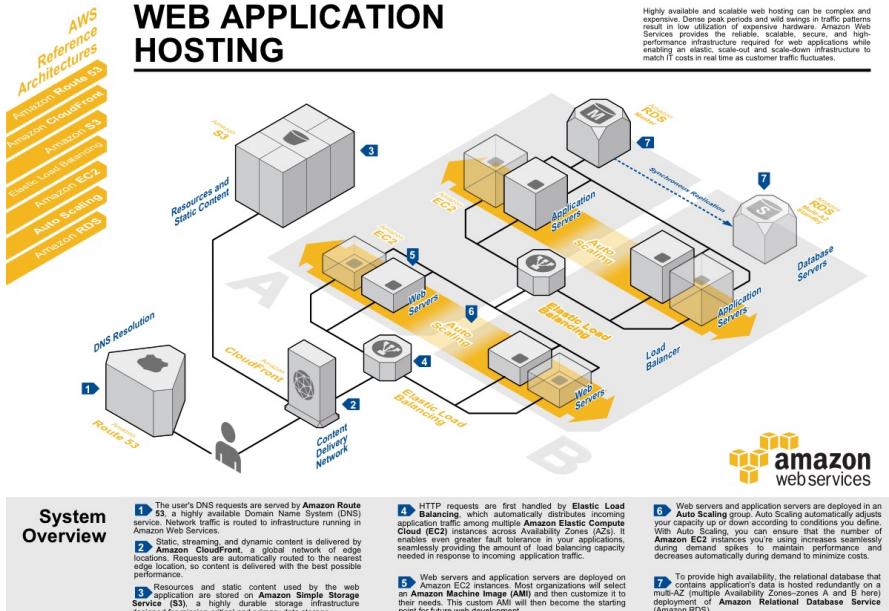
Decoupling of Components

- Enables independent scalability
- Fault Tolerance and High Availability
- Four important mechanisms to **decouple**
 - Load Balancers
 - Message Queues
 - Message Topics
 - Service Registries

Review

Architectural Patterns

Review



<https://aws.amazon.com/architecture>

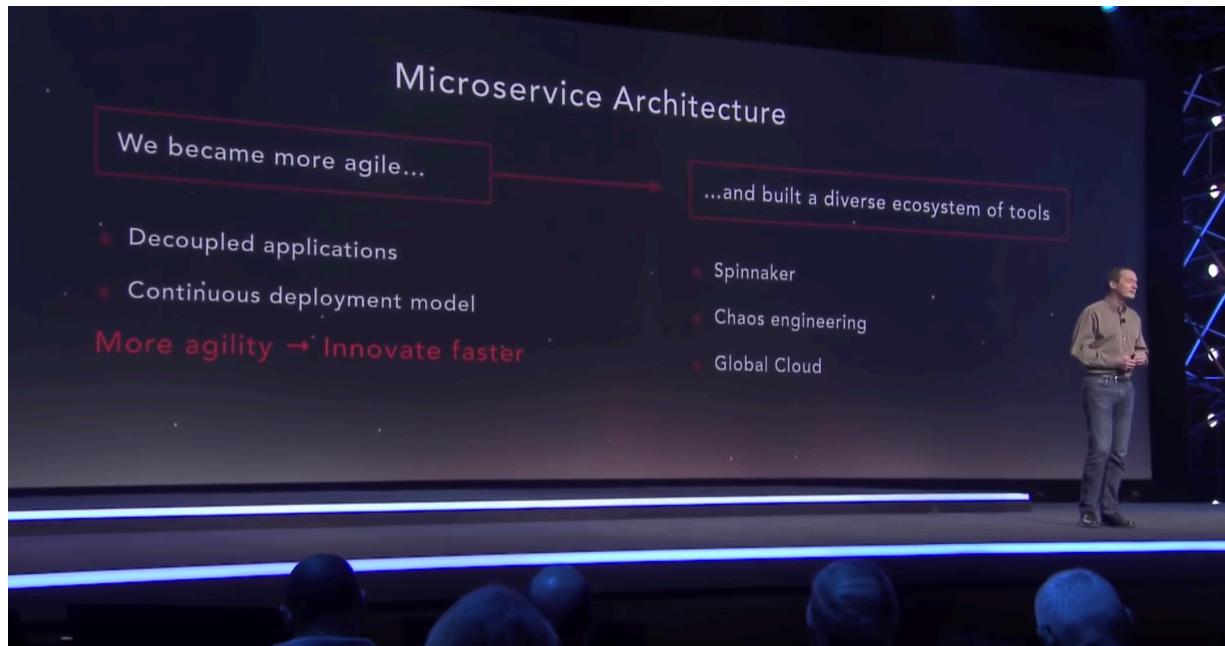
What is Cloud Native?

- Container based
 - Capabilities
- Elastic
- Microservices
 - Polyglot
- DevOps



AWS Biggest Customer?

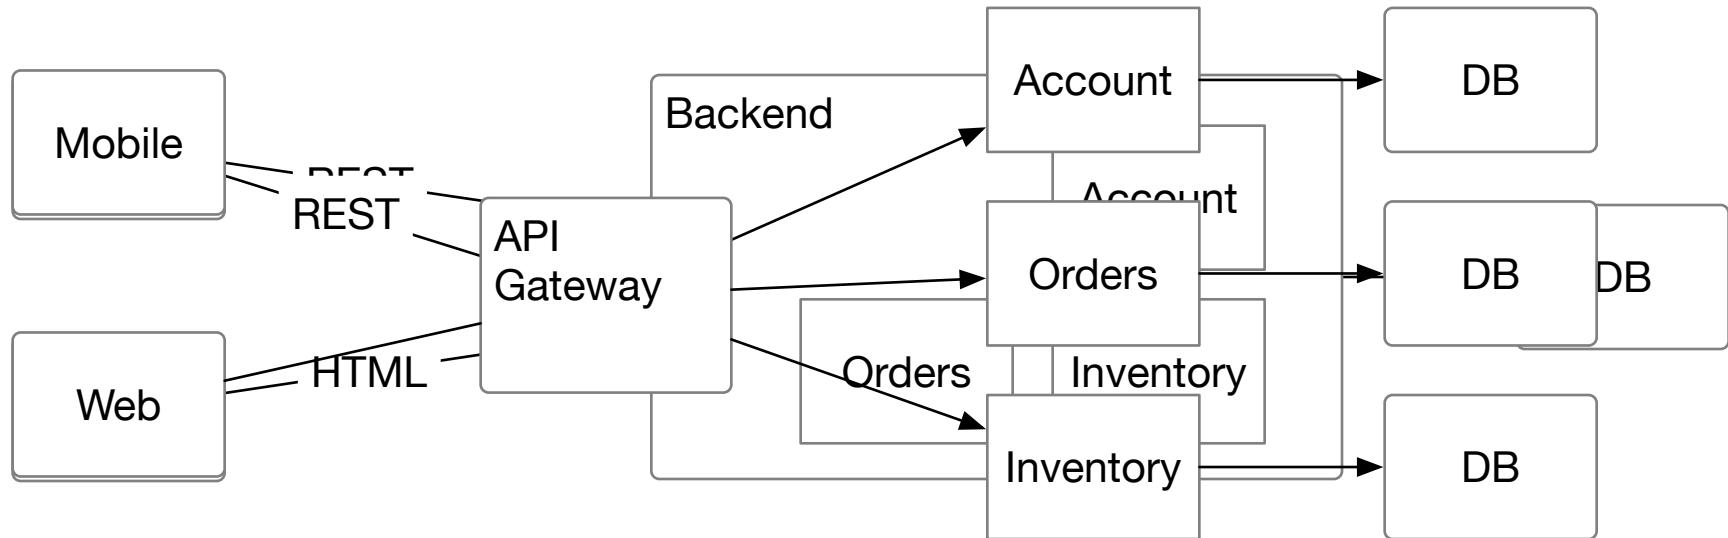
- Netflix
- >100.000 instances
- +500 microservices
- 1 application change per second



What Are Microservices?

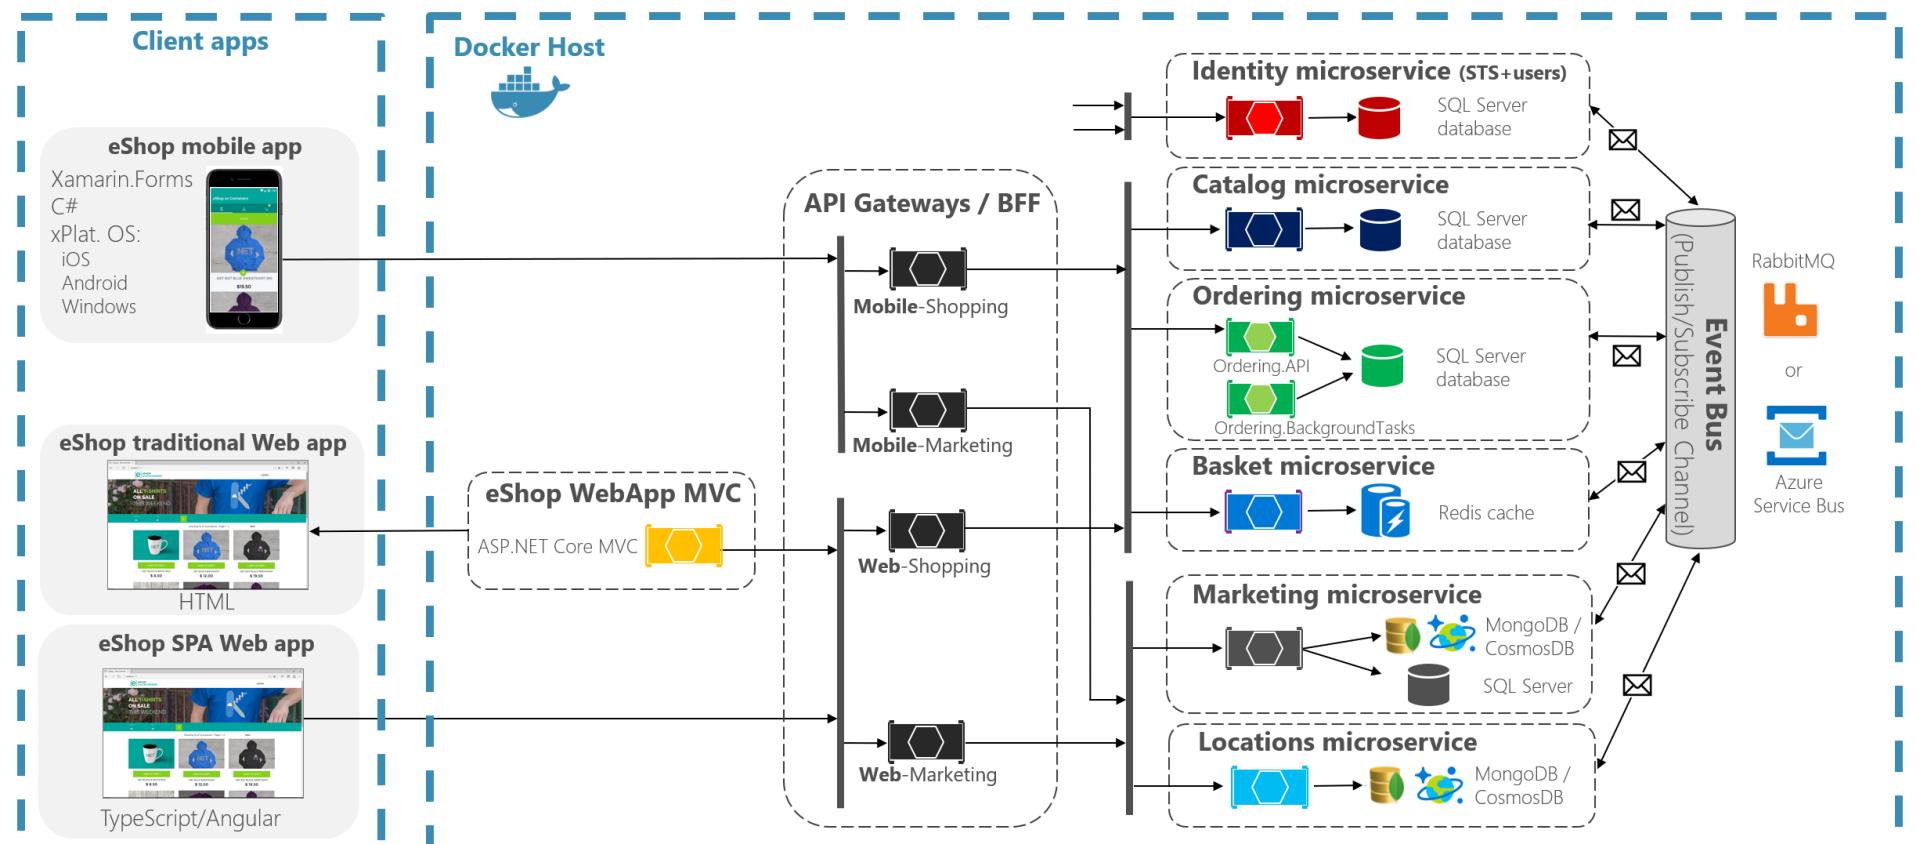
- Martin Fowler first used the term 2014
- Opposite “monolithic” applications
- Organisational agility
 - Rolling releases
 - Hot swapping
 - No down time
- Through
 - Horizontal Auto-Scaling
 - Design for Failure
 - Replication, Health Monitoring
 - Modularity
 - Decoupling
 - Containerised
 - API contracts
 - Discovery
 - Automation
 - DevOps
 - Observability

Breaking up the monolith



eShopOnContainers reference application

(Development environment architecture)

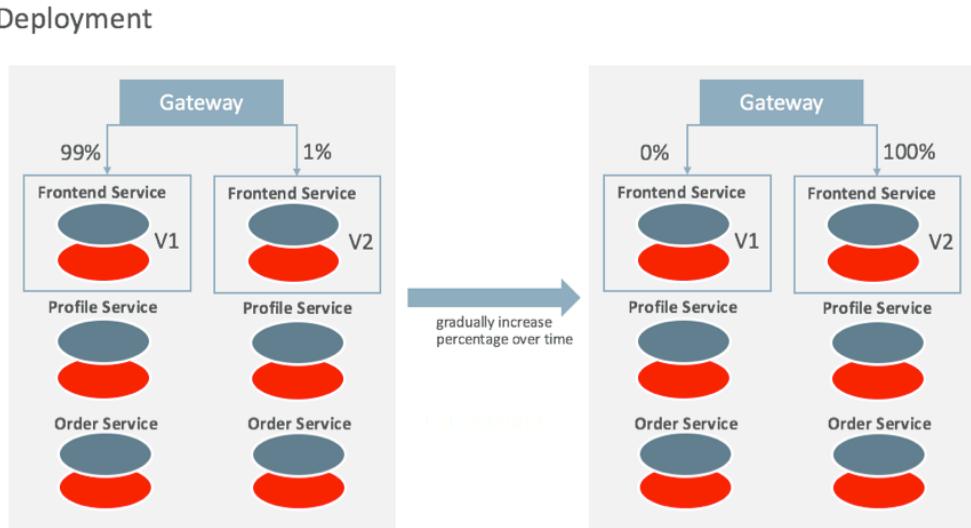


Upsides

- Strong module boundaries — great for large teams
- Independent deployment
 - Simple services are easier to deploy, and since they are autonomous, are less likely to cause system failures when they go wrong.
- Technology Diversity:
 - With microservices you can mix multiple languages, development frameworks and data-storage technologies.
- Re-use services

Service deployability

- Canary
 - a small percentage of users
- Blue/Green
 - Standby prod env
- A/B
 - Empirically evaluate a design



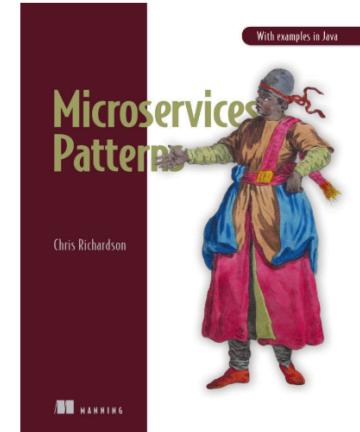
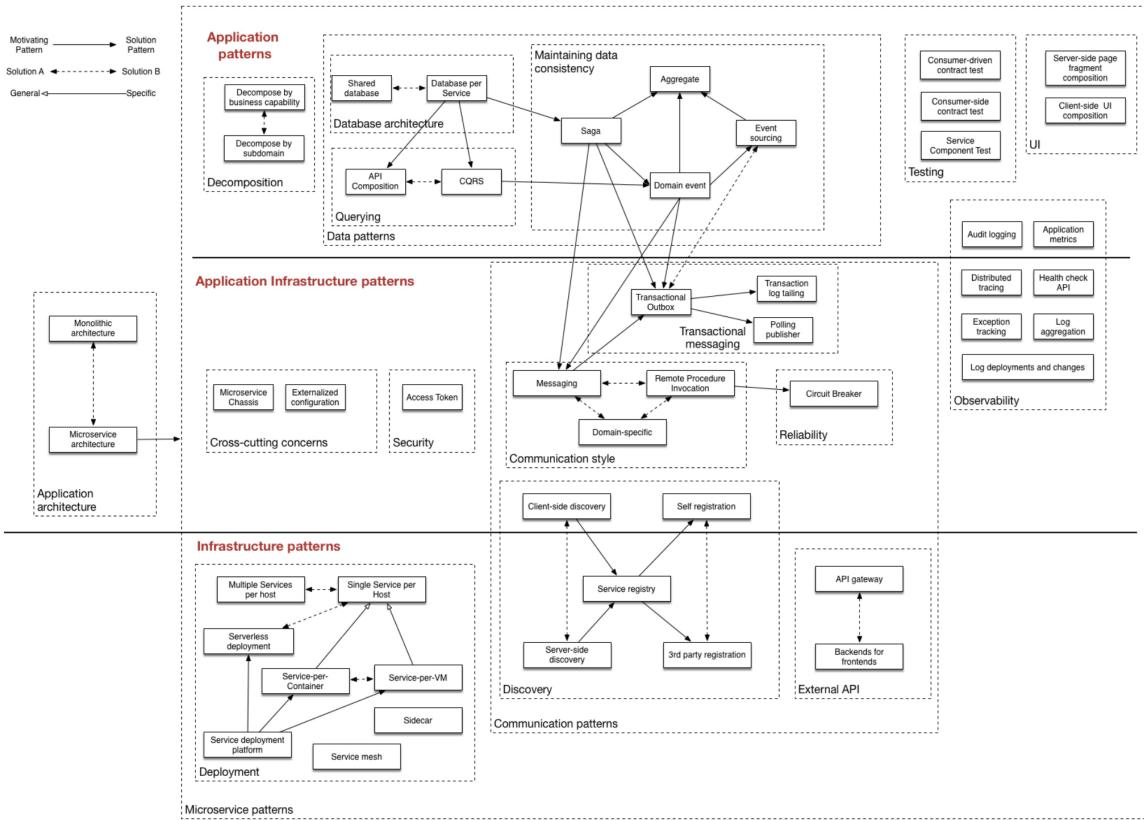
<https://blogs.oracle.com/developers/getting-started-with-microservices-part-four>

Downsides

- Distribution:
 - Distributed systems are harder to program, since remote calls are slow and are always at risk of failure.
- Eventual Consistency:
 - Maintaining strong consistency is extremely difficult for a distributed system, which means everyone has to manage eventual consistency.
- Operational Complexity:
 - You need a mature operations team to manage lots of services, which are being redeployed regularly.

Patterns and Strategies

- Chris Richardson
- microservices.io

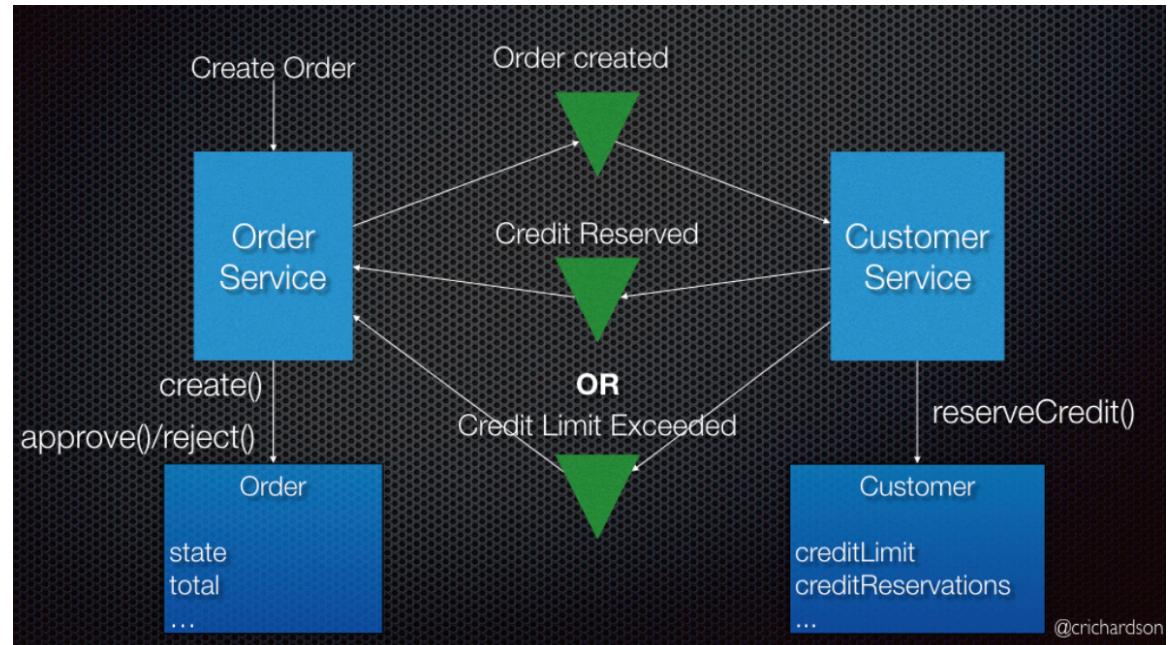


Data locality

- Each micro service has own DB - Limit access to peer service data to API only
- Org level: Allows each service to manage data in the way that makes the most sense for that service
- Performance benefit: caching data locally in microservices enables faster response
- Challenges:
 - invariants that span multiple services; query data that is owned by multiple services
 - Implementing business transactions that span multiple services is not straightforward
 - Avoided Distributed transactions (-> CAP theorem lecture)

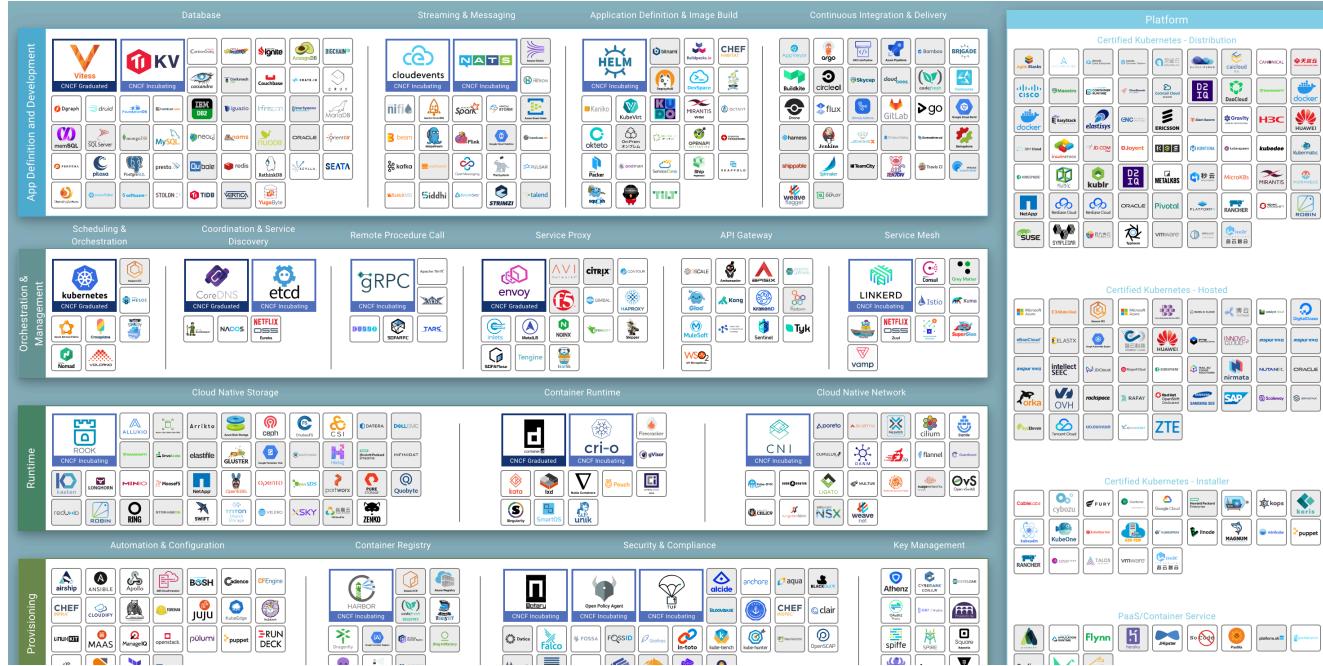
Eventual consistency

- “no shared database”
—> share state through HTTP
- No guarantee that those updates occur immediately
- Guarantee that they occur eventually



Saga Pattern -- <https://microservices.io/patterns/data/saga.html>

CNCF Landscape



CNCF Project Clusters

- [Database](#)
- [Streaming & Messaging](#)
- [Observability and Analysis](#)
 - [Monitoring](#)
 - [Logging](#)
 - [Tracing](#)
 - [Chaos Engineering](#)
- [Continuous Integration & Delivery](#)
- [Scheduling & Orchestration](#)
- [Coordination & Service Discovery](#)
- [Remote Procedure Call](#)
- [API Gateway](#)
- [Service Mesh](#)
- [Cloud Native Storage](#)
- [Container Runtime](#)
- [Cloud Native Network](#)
- [Automation & Configuration](#)
- [Container Registry](#)
- [Security & Compliance](#)
- [Key Management](#)
- [Service Proxy](#)

CNCF – Graduated projects

- Kubernetes
- Prometheus (9.8.2018)
 - monitoring system and time series database
- Envoy
 - Service Network Proxy
- CoreDNS
 - DNS server and Service Discovery
- Containerd (28.2.2019)
 - Container Runtime
- Fluentd (11.4.2019)
 - Logging
- Jaeger (4.11.2019)
 - Monitoring and tracing
- Vitess (6.11.2019)
 - Database
 - Sharded MySQL

Observability

- What is going, was and will happen
- "*Observability is a measure of how well internal states of a system can be inferred from knowledge of its external outputs*"
- Logging
- Monitoring & metrics
- Diagnostics
- Health checks
- Debugging
- Fault injection

Logging

- Local log files
 - Grows potentially unlimited (log rotate)
 - Ephemeral storage
- Aggregated Logs
 - Syslog
 - Central server
 - ElasticSearch, Logstash, Kibana



- ELK in Kubernetes
- side car logstash container
- Application level or platform level
- Fluentd



fluentd

Metrics

- More than logs
 - CPU, load average, interrupts
 - Free memory
 - Network packets received, transmitted, dropped
 - Disk space
- Identify bottlenecks
- Monitor the utilisation
- Identify hardware failures
- diagnosing



Prometheus

- TSDB
 - Efficient storage and indexing
 - Very compressable
- Prometheus
- Kubernetes metrics-server
 - Runs as a pod
 - Connects to the kubelet to retrieve pod and node cgroup information.
 - Exposes /metrics endpoint to be scraped by Prometheus

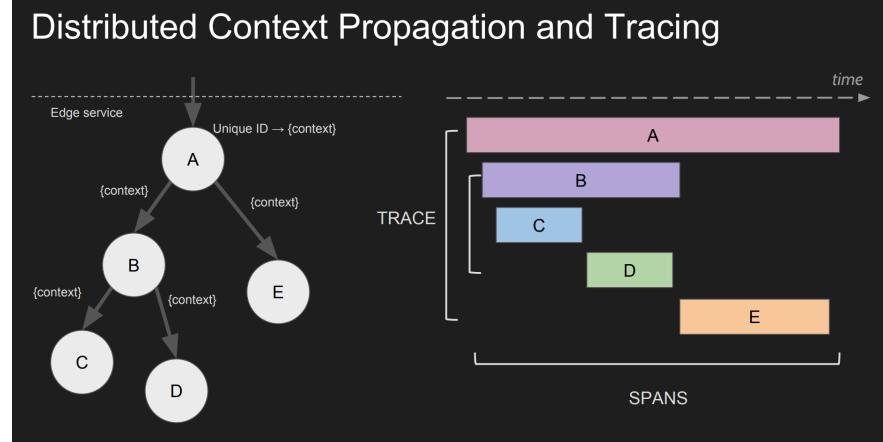
Jaeger



- Open sourced by Uber
- Logs contain messages from many requests -> hard to tell a story
- [OpenTracing](#)
- Trace single request across your microservices
- Context per request, propagate across call boundaries

```
io.opentracing.Tracer tracer = ...;
```

```
Span span = tracer.buildSpan("someWork").start();
```

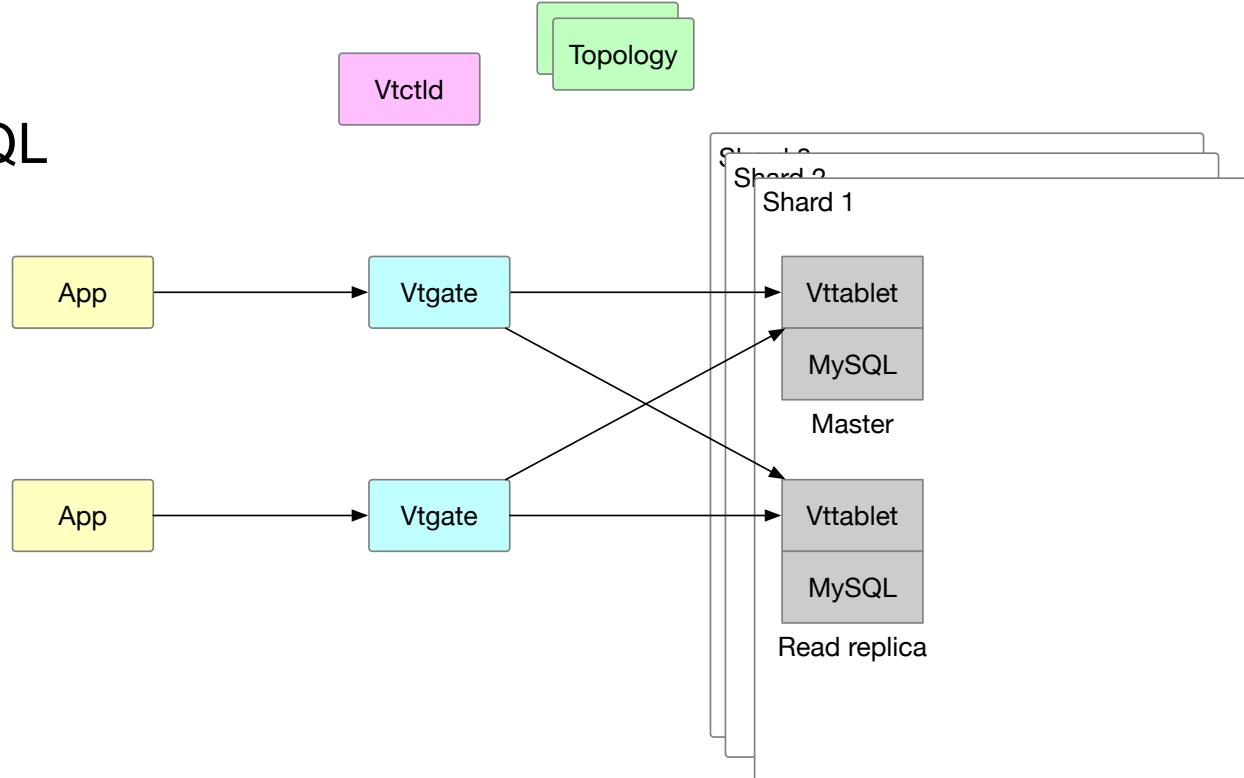


- Aggregated logs & profiling per request and process

Vitesse

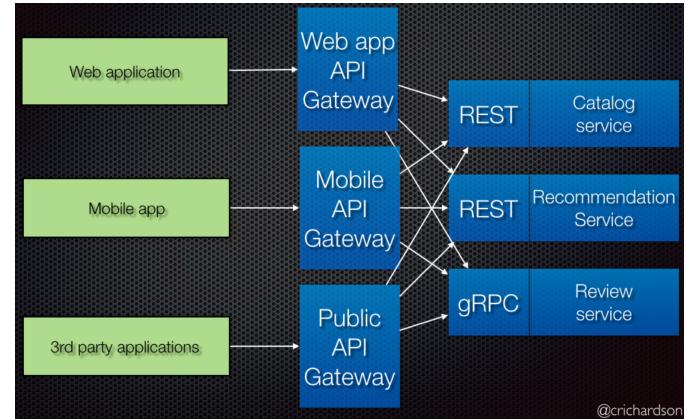


- Sharding for MySQL
- HA
- Pinterest, Slack, Github

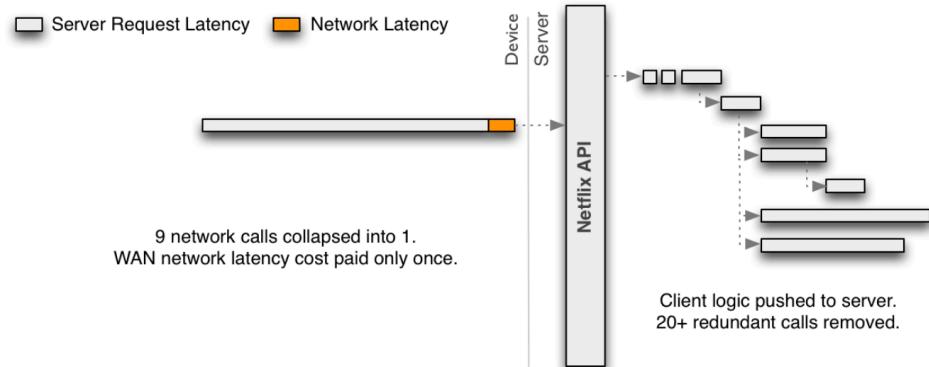


API Gateway

- Manage granularity of APIs
- Hide services from clients
- Hide diverse set of protocols
- Open API Spec
 - Integrated design, documentation, and implementation
- Manage network latency



<https://microservices.io/patterns/apigateway.html>



Chaos Monkey

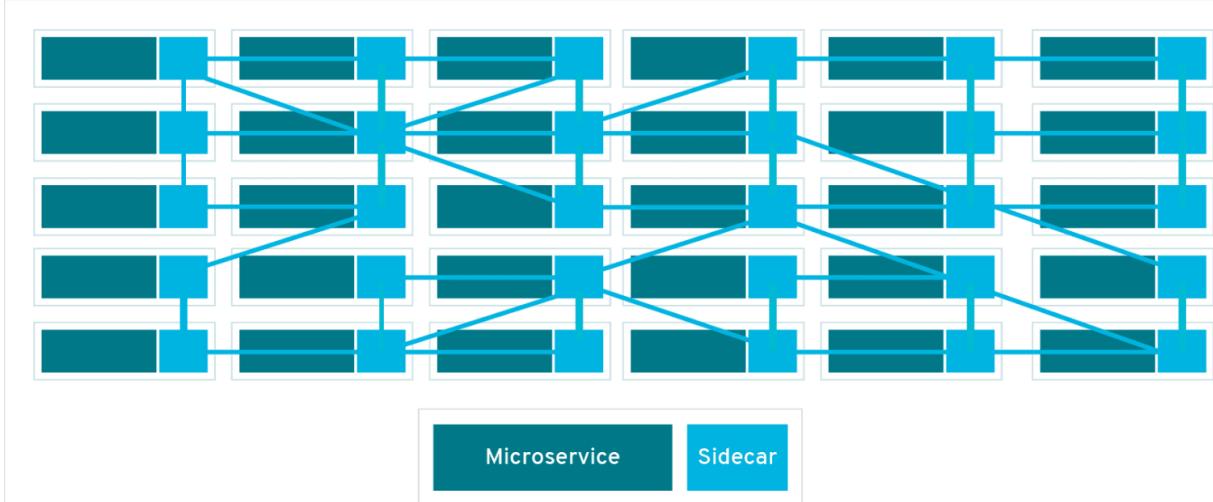
- For three days in August 2008, Netflix couldn't ship DVDs to customers because of a major database corruption.
 - Moved to a more resilient architecture
- Chaos Monkey randomly takes virtual machines
- Netflix closed all their DCs
- Netflix keeps backups of everything in [Google Cloud Storage](#)



LINKERD

Service Meshes

- Networking model
- requests are routed between microservices through proxies in their own infrastructure layer.
- Sidecar pattern ([Burns et al.](#))
- Security and Identity
- Routing Management
- Proxy collects data
 - Performance
 - Locates problem pods
 - Applies patterns
 - circuit-breaking,
 - latency-aware load balancing,
 - eventually consistent (“advisory”)
 - service discovery,
 - retries, and deadlines



Thanks!